

# Fundamentals of Accelerating Computing with CUDA C/C++

NVIDIA Deep Learning Institute

## About Me

Michel Agustin Herrera Sanchez, MSc, CQF

Master in High Performance Computing, University of A Coruña, Spain

OpenACC Standard Certified HPC GPU Mentor

NVIDIA DLI Certified Instructor – Accelerated Computing

BNP Paribas, Risk Systems Developer

Berlin, Germany

# Workshop Outline

## **Introduction (15 min)**

Accelerating Applications with CUDA C/C++ (120min)

Break (15 min)

Managed Accelerated Application Memory with CUDA C/C++ (120min)

Break (15 min)

Asynchronous Steaming and Visual Profiling for Accelerated Applications with CUDA/C++ (120min)

Final Review (15 min)

CUDA C is essentially C/C++ with a few extensions that allow one to execute functions on the GPU using many threads in parallel.

# Parallel Architecture – NVIDIA Tesla V100 GPU



7.8 TFLOPS FP64  
15.7 TFLOPS FP32  
900GB/s Memory Bandwidth

84 SMs  
5376 FP32 cores  
5376 INT32 cores  
2688 FP64 cores  
672 Tensor cores

Each SM has:  
64 FP32 cores  
64 INT32 cores  
32 FP64 cores  
8 Tensor Cores  
4 Texture units

Memory Size 16GB  
L2 Cache Size 6144KB  
Shared Memory Size/SM 96KB  
Register File Size/SM 256KB

# Parallel Architecture – Volta V100 Streaming Multi Processor (SM)



# GPU Computing

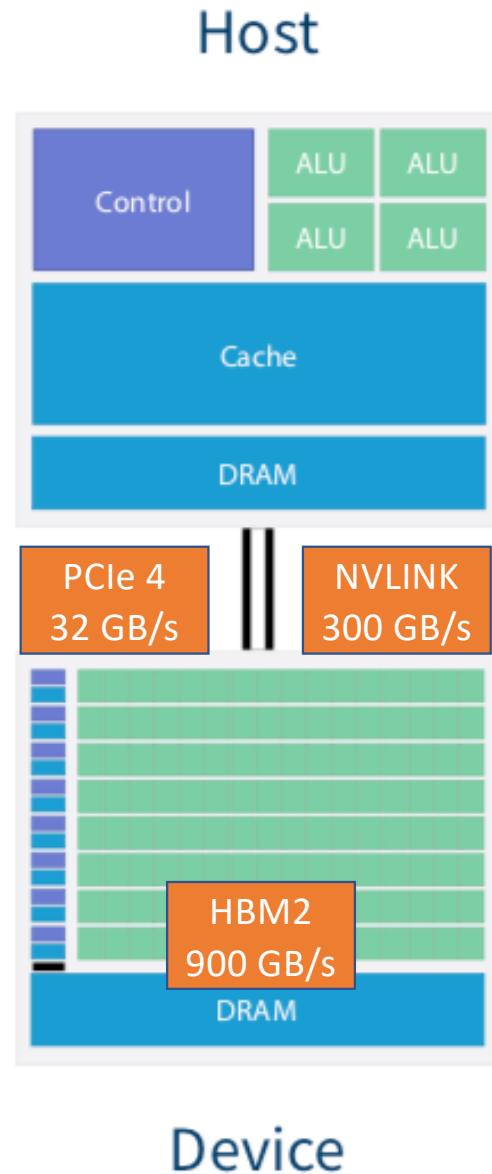
## Differences between CPU and GPU Architectures

### CPU

- More suitable for low latency workloads
- CPU has bigger memory caches
- Higher Processor Clock Rate
- Execution Model SIMD/SMT
- Data doesn't need to be moved from Host to Device to be processed

### GPU

- More suitable for high throughput workloads
- GPU has less memory capacity but higher main memory bandwidth
- Lower Processor Clock Rate
- Execution Model SIMT
- Data Transfer between host and device is a Limiting performance factor

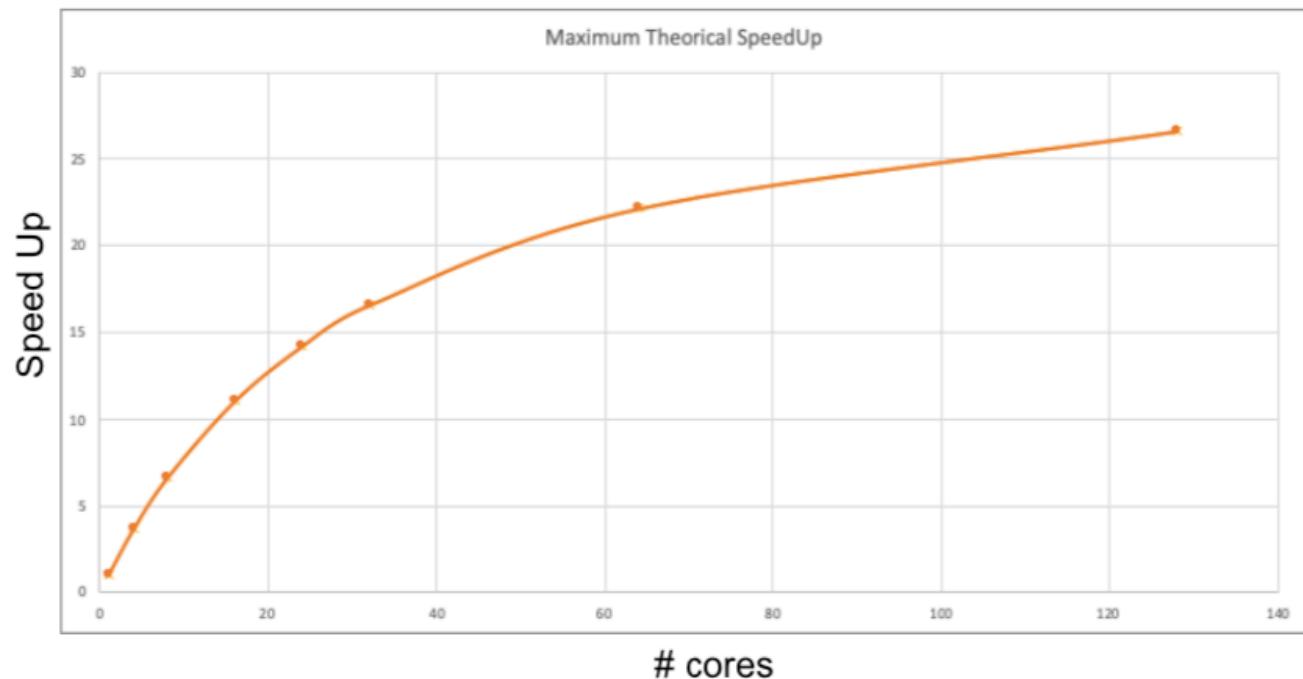


# Theoretical Speed Up using Amdahls Law

Time(%)	Total Time (ns)	Instances	Average (ns)	Minimum (ns)	Maximum (ns)	StdDev (ns)	NVTX
33.9	18,600,340,372	1	18,600,340,372.0	18,600,340,372	18,600,340,372	0.0	total
33.1	18,136,889,464	1	18,136,889,464.0	18,136,889,464	18,136,889,464	0.0	main_simulation_loop
20.6	11,304,100,596	5,000	2,260,820.1	2,175,354	9,488,915	137,055.8	jacobi
12.4	6,822,128,866	5,000	1,364,425.8	1,338,596	6,271,088	140,306.2	swap
0.0	8,352,028	1	8,352,028.0	8,352,028	8,352,028	0.0	init
0.0	1,970,085	5,000	394.0	136	746,645	10,559.8	calculate_error

$$S(n) = \frac{1}{(1 - P) + \frac{P}{n}}$$

# cores [n]	Theoretical SpeedUp [S(n)]
1	1
4	3.669724771
8	6.611570248
16	11.03448276
24	14.20118343
32	16.58031088
64	22.14532872
128	26.61122661
Parallizable Fraction[P]	0.97



# Application Performance

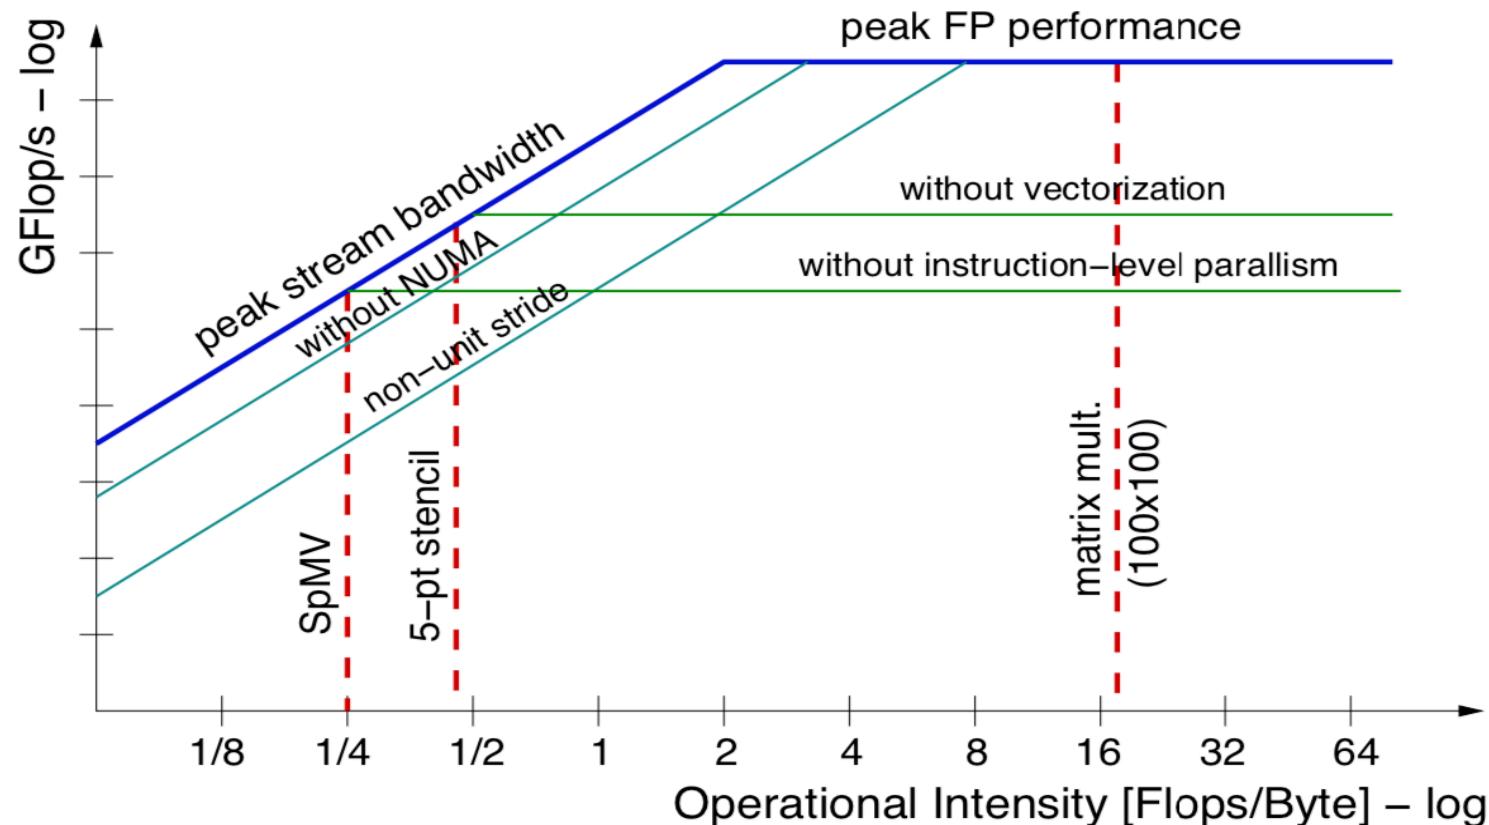
## Memory Bound vs Compute Bound

- Compute to global memory access ratio (Arithmetic Intensity) is defined as the number of FLOPs performed for each byte access from the global memory within a kernel.
- Memory Bound Performance
  - Increase Speed up by making an efficient/reducing memory access
- Compute Bound Performance
  - Increase Speed up by reducing/optimizing compute operations

$$\frac{\#ops}{\#bytes} < \frac{BW_{math}}{BW_{mem}}$$

$$\frac{\#ops}{\#bytes} > \frac{BW_{math}}{BW_{mem}}$$

# The Roof Line Model



Williams, Waterman, Patterson 2008

# Programming GPUs with CUDA C/C++

# Programming GPUs with CUDA C/C++

```
#include <iostream>
#include <math.h>

// function to add the elements of two arrays
void add(int n, float *x, float *y) {
    for (int i = 0; i < n; i++)
        y[i] = x[i] + y[i];
}

int main(void) {
    int N = 1<<20; // 1M elements
    float *x = (float*) malloc( N*sizeof(float) );
    float *y = (float*) malloc( N*sizeof(float) );

    // Initialize x, y data
    for (int i = 0; i < N; i++) {
        x[i] = 1.0f; y[i] = 2.0f;
    }
}

// Run kernel on 1M elements on the CPU
add(N, x, y);

// Check for errors (all values should be 3.0f)
float maxError = 0.0f;

for (int i = 0; i < N; i++)
    maxError = fmax(maxError, fabs(y[i]-3.0f));

std::cout << "Max error: " << maxError << std::endl;

// Free memory
free(x);
free(y);

return 0;
```

# “add” CUDA Kernel (Monolithic) implementation

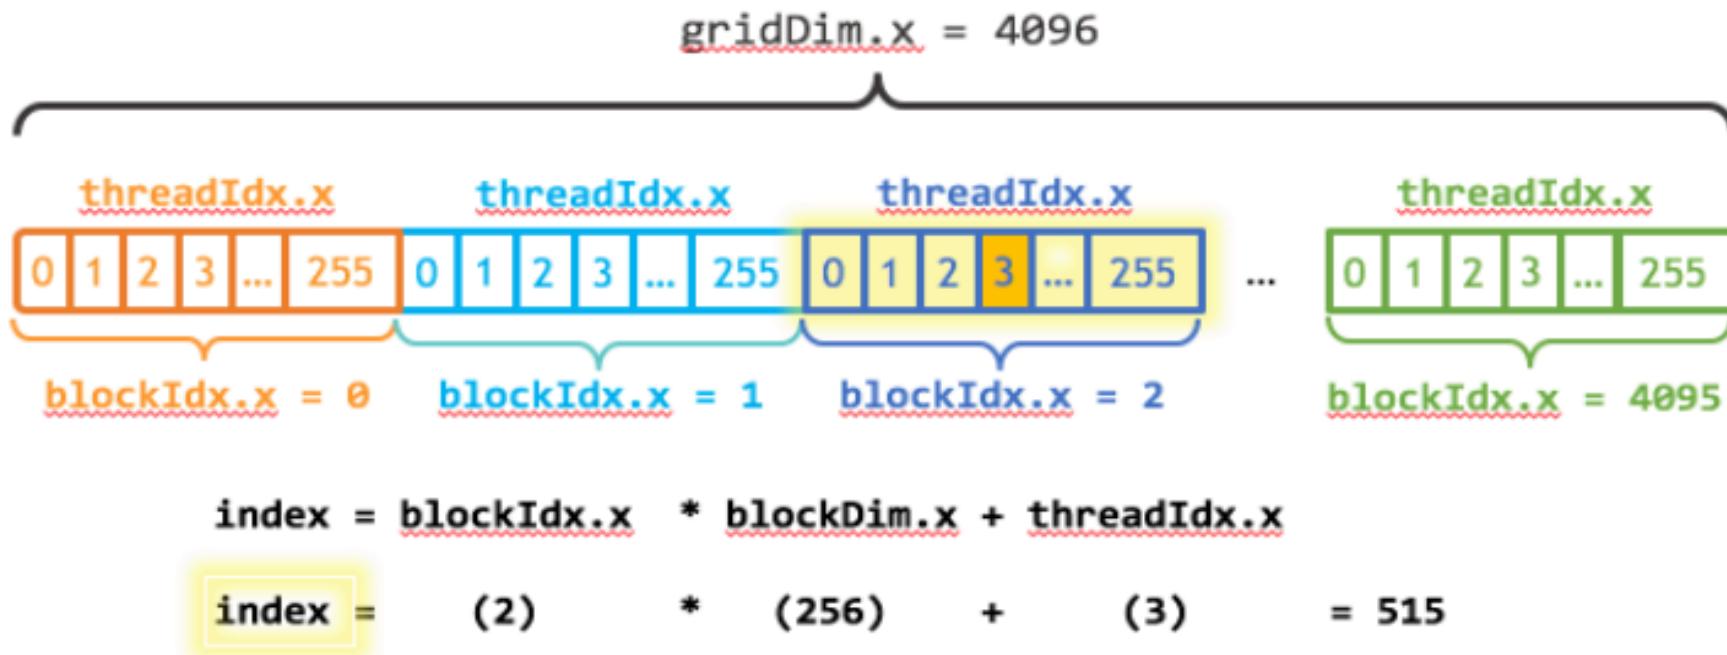
```
void add(int n, float *x, float *y) {  
  
    for (int i = 0; i < n; i++)  
        y[i] = x[i] + y[i];  
}
```

Parallelize a loop using a kernel.  
Launch One Thread Per Data Element

	0	1	2	3	4	5	6	7	8	9	10	11
x	3	7	0	1	4	0	0	4	5	3	1	0
y	2	4	2	1	8	3	9	5	5	1	2	1
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
result	5	11	2	2	12	3	9	9	10	4	3	1

```
__global__ void add(int n, float *x, float *y) {  
  
    int index = blockIdx.x * blockDim.x + threadIdx.x;  
  
    if (index < n) {  
        y[index] = x[index] + y[index];  
    }  
}
```

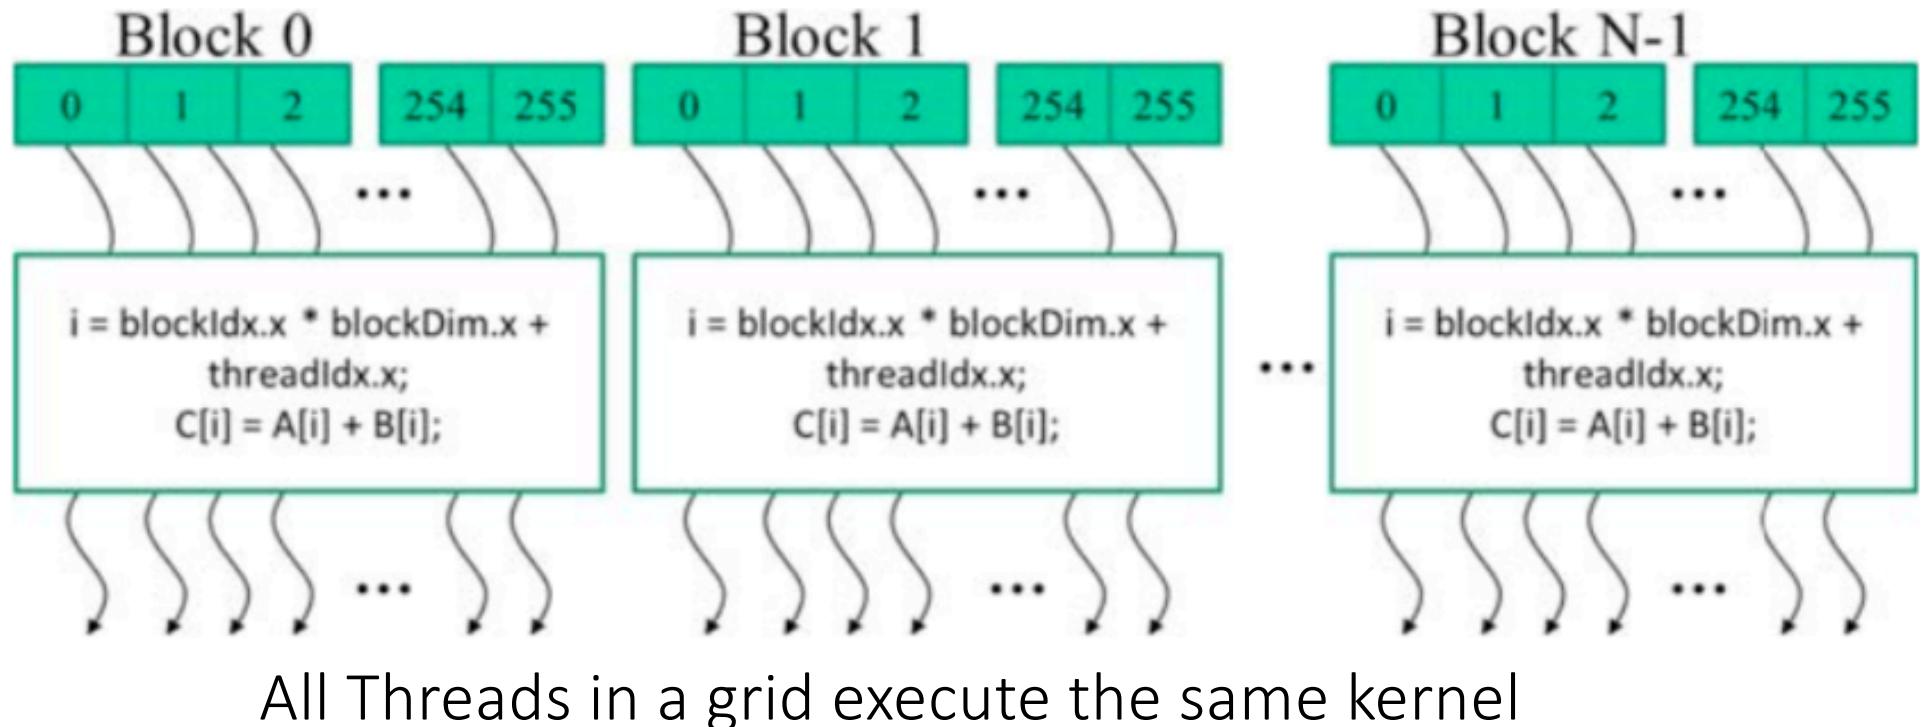
# Grid / Thread Blocks / Threads



# CUDA Kernel Execution Configuration

```
int blockSize = 256;  
int numBlocks = (N + blockSize - 1) / blockSize;  
add<<<numBlocks, blockSize>>>(N, x, y);
```

# Grid / Thread Blocks / Threads

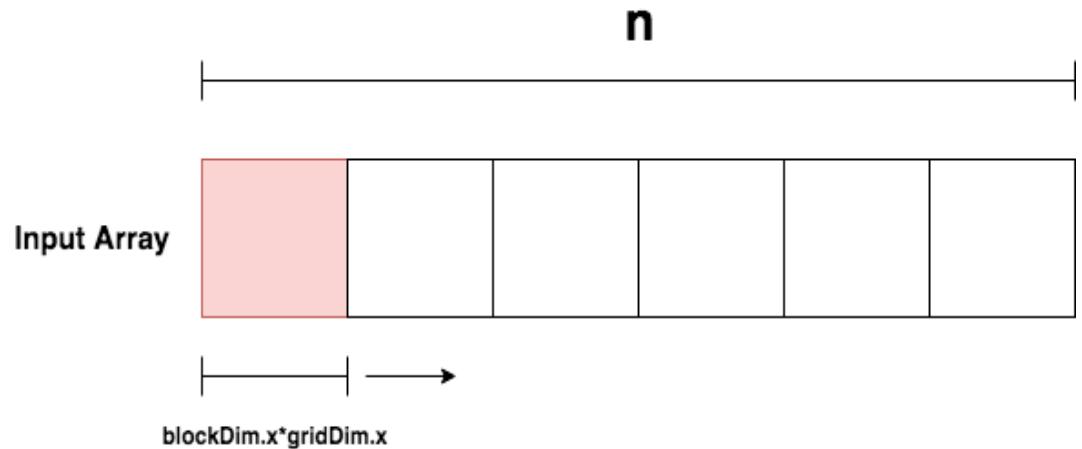


<https://docs.nvidia.com/cuda/#programming-guides>

Wen Mei Hwu, David B Kirk Programming Massively Parallel Processors A Hands-on Approach 4th 2022

# “add” CUDA Kernel (Grid Stride Loop)

```
__global__ void add(int n, float *x, float *y) {  
  
    int index = blockIdx.x * blockDim.x + threadIdx.x;  
  
    int stride = blockDim.x * gridDim.x;  
  
    for (int i = index; i < n; i += stride)  
        y[i] = x[i] + y[i];  
}
```



<https://developer.nvidia.com/blog/cuda-pro-tip-write-flexible-kernels-grid-stride-loops/>

<https://alexminnaar.com/2019/08/02/grid-stride-loops.html>

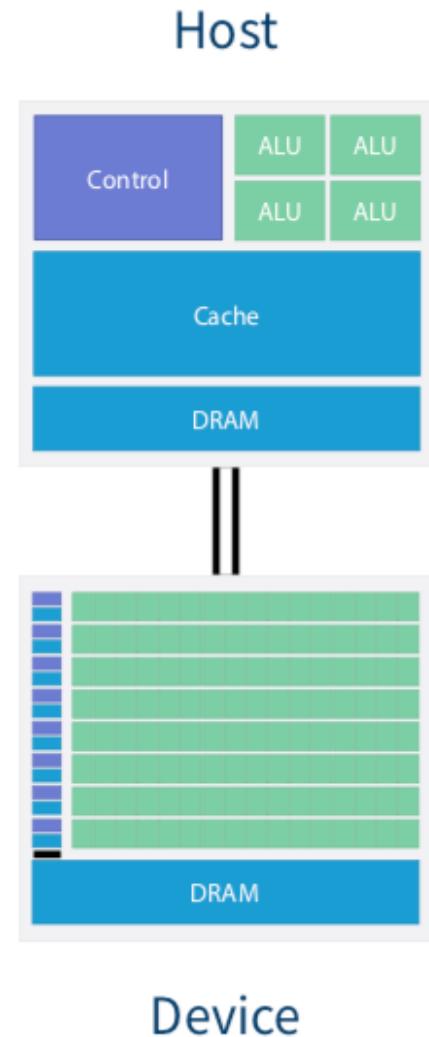
[https://www.nvidia.com/content/GTC-2010/pdfs/2238\\_GTC2010.pdf](https://www.nvidia.com/content/GTC-2010/pdfs/2238_GTC2010.pdf)

# Grid Stride Loop Execution Configuration

```
int numSMs;  
cudaDeviceGetAttribute(&numSMs, cudaDevAttrMultiProcessorCount, 0);  
  
int blockSize = 256;  
int numBlocks = 32 * numSMs;  
add<<<numBlocks, blockSize>>>(N, x, y);
```

# Host/Device Memory Model

```
float *x, *y, *x_d, *y_d;  
  
int size = N*sizeof(float);  
  
// Allocate Host Memory  
x = (float*) malloc( size );  
y = (float*) malloc( size );  
  
// Allocate Device Memory  
cudaMalloc((void**)&x_d, size);  
cudaMalloc((void**)&y_d, size);  
  
// Initialize x, y data arrays ....  
  
// Copy Data from Host to Device  
cudaMemcpy(x_d, x, size, cudaMemcpyHostToDevice);  
cudaMemcpy(y_d, y, size, cudaMemcpyHostToDevice);  
  
// Execute CUDA Kernel ....  
// ...  
  
// Copy Data from Device to Host  
cudaMemcpy(y, y_d, size, cudaMemcpyDeviceToHost);  
  
// Free memory  
cudaFree(x_d); cudaFree(y_d);  
free(x); free(y);
```

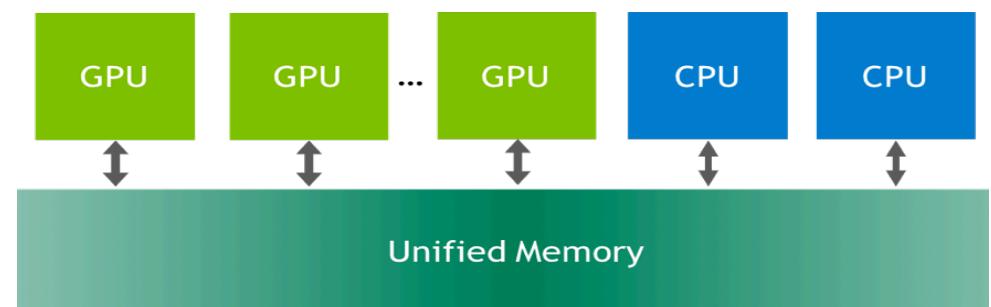


# CUDA Unified Memory Model

```
// Allocate Memory accessible from CPU or GPU
```

```
float *x, *y;  
cudaMallocManaged(&x, N*sizeof(float));  
cudaMallocManaged(&y, N*sizeof(float));
```

```
...  
// Free memory  
cudaFree(x);  
cudaFree(y);
```



# Programming GPUs with CUDA C/C++

```
#include <iostream>
#include <math.h>

// C CUDA Kernel
__global__ void add(int n, float *x, float *y) { //...}

int main(void) {
    int N = 1<<20; // 1M elements
    float *x, *y, *x_d, *y_d;
    int size = N*sizeof(float);

    // Allocate Host Memory
    x = (float*) malloc( size );
    y = (float*) malloc( size );

    // Allocate Device Memory
    cudaMalloc((void**)&x_d, size);
    cudaMalloc((void**)&y_d, size);

    // Initialize data x, y on host ....

    // Copy Data from Host to Device
    cudaMemcpy(x_d, x, size, cudaMemcpyHostToDevice);
    cudaMemcpy(y_d, y, size, cudaMemcpyHostToDevice);
}

// Run kernel on 1M elements on the GPU
int blockSize = 256;
int numBlocks = (N + blockSize - 1) / blockSize;
add<<<numBlocks, blockSize>>>(N, x, y);

// Copy Data from Device to Host
cudaMemcpy(y, y_d, size, cudaMemcpyDeviceToHost);

// Check for errors (all values should be 3.0f)
float maxError = 0.0f;

for (int i = 0; i < N; i++)
    maxError = fmax(maxError, fabs(y[i]-3.0f));

std::cout << "Max error: " << maxError << std::endl;

// Free memory
cudaFree(x_d);  cudaFree(y_d);

free(x); free(y);

return 0;
}
```

# Programming GPUs with CUDA C/C++

```
#include <iostream>
#include <math.h>

// C CUDA Kernel
__global__ void add(int n, float *x, float *y) { //...
}

int main(void) {
    int N = 1<<20; // 1M elements
    float *x, *y;

    // Allocate Unified Memory
    cudaMallocManaged(&x, N*sizeof(float));
    cudaMallocManaged(&y, N*sizeof(float));

    // Initialize x, y data on host
    for (int i = 0; i < N; i++) {
        x[i] = 1.0f; y[i] = 2.0f;
    }
}

// Run kernel on 1M elements on the GPU
int blockSize = 256;
int numBlocks = (N + blockSize - 1) / blockSize;
add<<<numBlocks, blockSize>>>(N, x, y);

// Wait for GPU to finish before accessing on host
cudaDeviceSynchronize();

// Check for errors (all values should be 3.0f)
float maxError = 0.0f;

for (int i = 0; i < N; i++)
    maxError = fmax(maxError, fabs(y[i]-3.0f));

std::cout << "Max error: " << maxError << std::endl;

// Free memory
cudaFree(x);
cudaFree(y);

return 0;
}
```

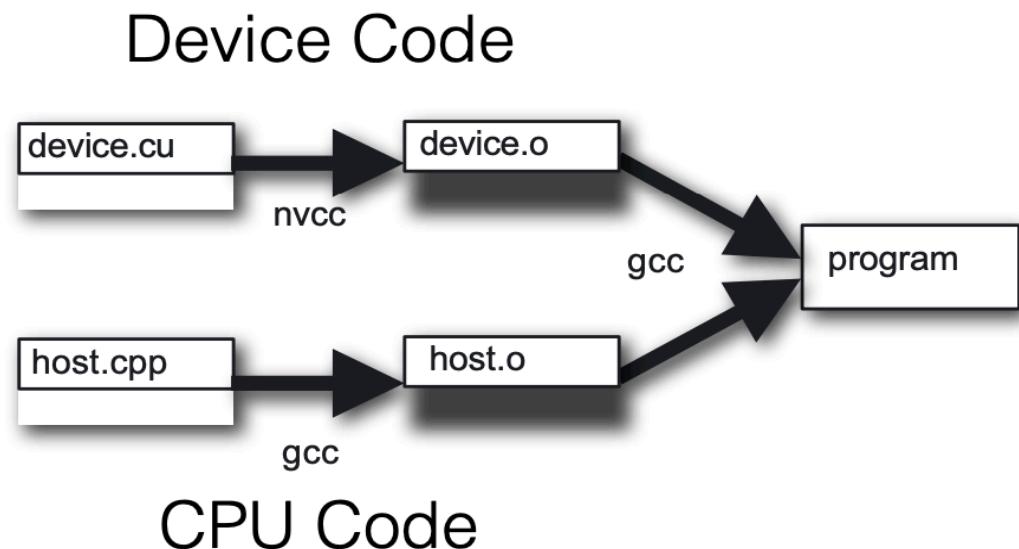
# Compiling, Executing and Profiling CUDA C/C++ Kernels

# GPU Compilation and Execution

```
> nvcc add.cu -o add_cuda
```

```
> ./add_cuda
```

```
Max error: 0.000000
```



# Profiling

```
> nsys profile --stats=true --force-overwrite=true -o nsys-report ./add_cuda  
Max error: 0.000000
```



Time(%)	Total Time (ns)	Instances	Average (ns)	Minimum (ns)	Maximum (ns)	StdDev (ns)	NVTX
38.4	3,282,267,380	1	3,282,267,380.0	3,282,267,380	3,282,267,380	0.0	total
33.0	2,824,520,551	1	2,824,520,551.0	2,824,520,551	2,824,520,551	0.0	main_simulation_loop
14.3	1,223,043,756	5,000	244,608.8	227,902	1,236,437	38,328.4	jacobi
14.2	1,210,402,922	5,000	242,080.6	226,149	477,625	33,842.8	swap
0.1	7,693,156	1	7,693,156.0	7,693,156	7,693,156	0.0	init
0.0	1,636,038	5,000	327.2	136	547,247	7,740.7	calculate_error

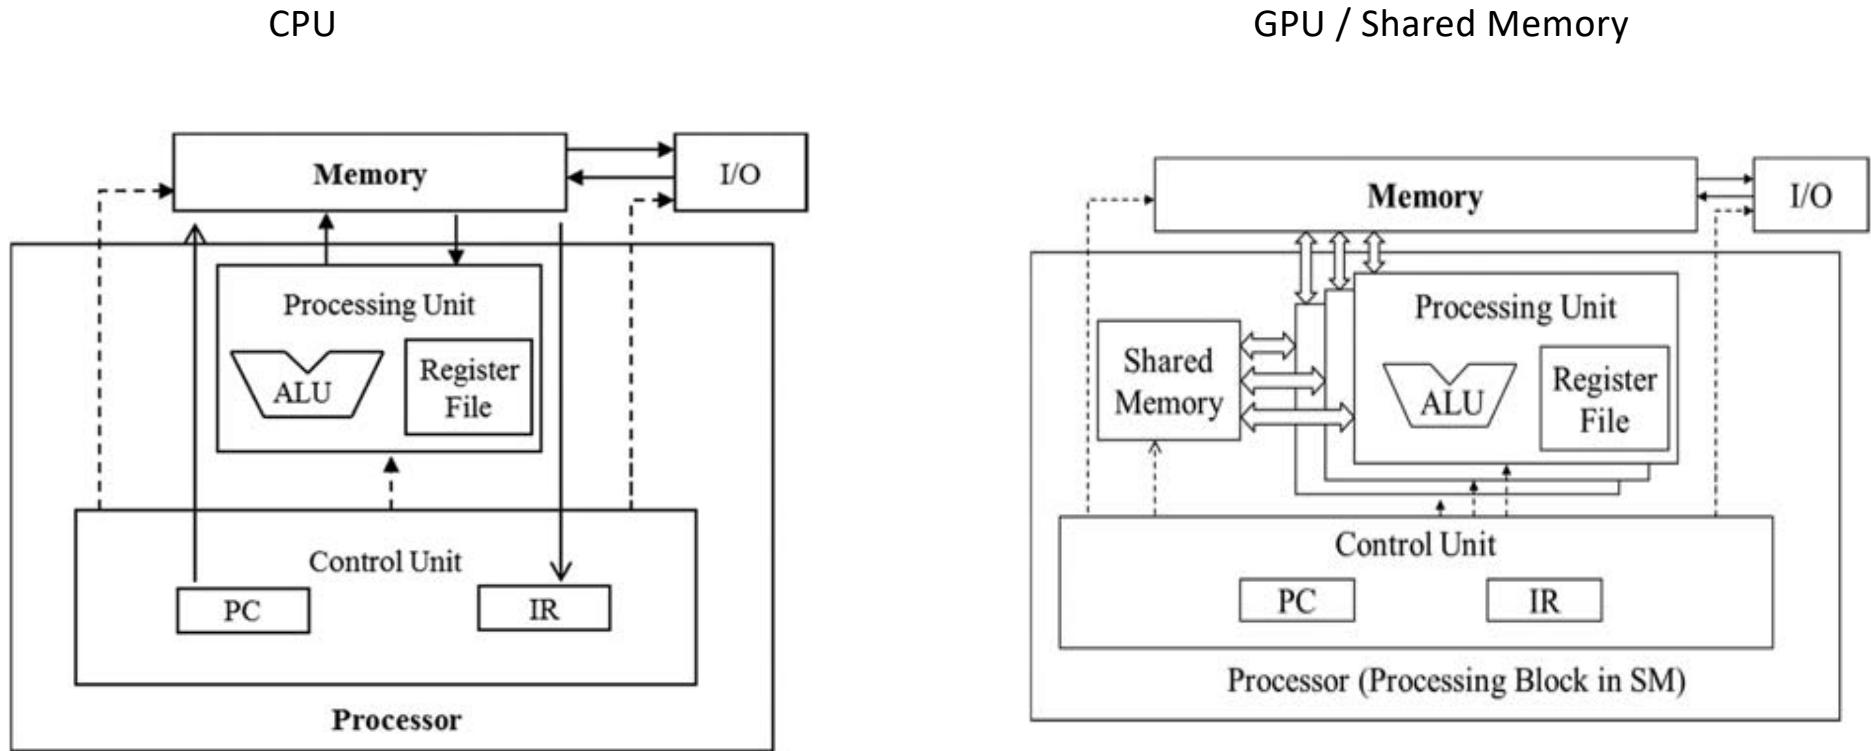
<https://developer.nvidia.com/blog/cuda-pro-tip-generate-custom-application-profile-timelines-nvtx/>

<https://docs.nvidia.com/nsight-systems/UserGuide/index.html>

# Performance Summary

	Laptop (GeForce GT 750M)	Server (Tesla K80)		
Version	Time	Bandwidth	Time	Bandwidth
1 CUDA Thread	411ms	30.6 MB/s	463ms	27.2 MB/s
1 CUDA Block	3.2ms	3.9 GB/s	2.7ms	4.7 GB/s
Many CUDA Blocks	0.68ms	18.5 GB/s	0.094ms	134 GB/s

<https://developer.nvidia.com/blog/even-easier-introduction-cuda/>



<https://www.olcf.ornl.gov/cuda-training-series/>

<https://www.olcf.ornl.gov/wp-content/uploads/2019/12/02-CUDA-Shared-Memory.pdf>

Wen Mei Hwu, David B Kirk Programming Massively Parallel Processors A Hands-on Approach 4th 2022

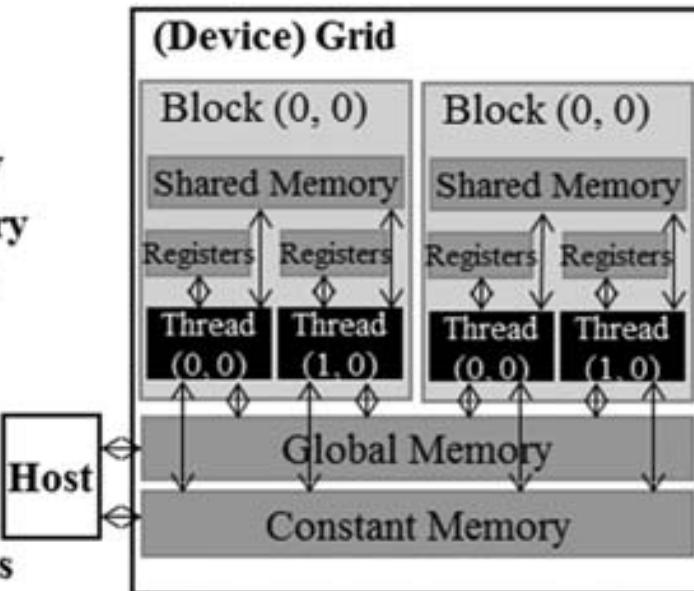
# Overview of CUDA Device Memory Model

Device code can:

- R/W per-thread **registers**
- R/W per-thread **local memory**
- R/W per-block **shared memory**
- R/W per-grid **global memory**
- Read only per-grid **constant memory**

Host code can

- Transfer data to/from per grid **global and constant memories**



Variable declaration	Memory	Scope	Lifetime
Automatic variables other than arrays	Register	Thread	Grid
Automatic array variables	Local	Thread	Grid
<code>__device__ __shared__ int SharedVar;</code>	Shared	Block	Grid
<code>__device__ int GlobalVar;</code>	Global	Grid	Application
<code>__device__ __constant__ int ConstVar;</code>	Constant	Grid	Application

# Matrix Matrix Multiplication

Access order				
thread <sub>0,0</sub>	M <sub>0,0</sub> * N <sub>0,0</sub>	M <sub>0,1</sub> * N <sub>1,0</sub>	M <sub>0,2</sub> * N <sub>2,0</sub>	M <sub>0,3</sub> * N <sub>3,0</sub>
thread <sub>0,1</sub>	M <sub>0,0</sub> * N <sub>0,1</sub>	M <sub>0,1</sub> * N <sub>1,1</sub>	M <sub>0,2</sub> * N <sub>2,1</sub>	M <sub>0,3</sub> * N <sub>3,1</sub>
thread <sub>1,0</sub>	M <sub>1,0</sub> * N <sub>0,0</sub>	M <sub>1,1</sub> * N <sub>1,0</sub>	M <sub>1,2</sub> * N <sub>2,0</sub>	M <sub>1,3</sub> * N <sub>3,0</sub>
thread <sub>1,1</sub>	M <sub>1,0</sub> * N <sub>0,1</sub>	M <sub>1,1</sub> * N <sub>1,1</sub>	M <sub>1,2</sub> * N <sub>2,1</sub>	M <sub>1,3</sub> * N <sub>3,1</sub>

FIGURE 5.6

Global memory accesses performed by threads in block<sub>0,0</sub>.

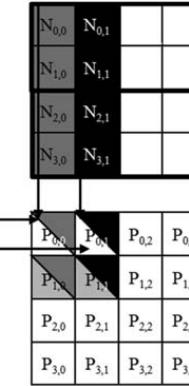


FIGURE 5.7

Tiling M and N to utilize shared memory.

	Phase 0		Phase 1		time	
	M <sub>0,0</sub>	N <sub>0,0</sub>	PValue <sub>0,0</sub> += Mds <sub>0,0</sub> *Nds <sub>0,0</sub> + Mds <sub>0,1</sub> *Nds <sub>1,0</sub>	M <sub>0,2</sub>		
thread <sub>0,0</sub>	M <sub>0,0</sub>	N <sub>0,0</sub>	PValue <sub>0,0</sub> += Mds <sub>0,0</sub> *Nds <sub>0,0</sub> + Mds <sub>0,1</sub> *Nds <sub>1,0</sub>	M <sub>0,2</sub>	N <sub>2,0</sub>	PValue <sub>0,0</sub> += Mds <sub>0,0</sub> *Nds <sub>0,0</sub> + Mds <sub>0,1</sub> *Nds <sub>1,0</sub>
thread <sub>0,1</sub>	M <sub>0,1</sub>	N <sub>0,1</sub>	PValue <sub>0,1</sub> += Mds <sub>0,0</sub> *Nds <sub>0,1</sub> + Mds <sub>0,1</sub> *Nds <sub>1,1</sub>	M <sub>0,3</sub>	N <sub>2,1</sub>	PValue <sub>0,1</sub> += Mds <sub>0,0</sub> *Nds <sub>0,1</sub> + Mds <sub>0,1</sub> *Nds <sub>1,1</sub>
thread <sub>1,0</sub>	M <sub>1,0</sub>	N <sub>1,0</sub>	PValue <sub>1,0</sub> += Mds <sub>1,0</sub> *Nds <sub>0,0</sub> + Mds <sub>1,1</sub> *Nds <sub>1,0</sub>	M <sub>1,2</sub>	N <sub>3,0</sub>	PValue <sub>1,0</sub> += Mds <sub>1,0</sub> *Nds <sub>0,0</sub> + Mds <sub>1,1</sub> *Nds <sub>1,0</sub>
thread <sub>1,1</sub>	M <sub>1,1</sub>	N <sub>1,1</sub>	PValue <sub>1,1</sub> += Mds <sub>1,0</sub> *Nds <sub>0,1</sub> + Mds <sub>1,1</sub> *Nds <sub>1,1</sub>	M <sub>1,3</sub>	N <sub>3,1</sub>	PValue <sub>1,1</sub> += Mds <sub>1,0</sub> *Nds <sub>0,1</sub> + Mds <sub>1,1</sub> *Nds <sub>1,1</sub>

FIGURE 5.8

Execution phases of a tiled matrix multiplication.

```

01      #define TILE_WIDTH 16
02      __global__ void matrixMulKernel(float* M, float* N, float* P, int Width) {
03
04          __shared__ float Mds[TILE_WIDTH][TILE_WIDTH];
05          __shared__ float Nds[TILE_WIDTH][TILE_WIDTH];
06
07          int bx = blockIdx.x;  int by = blockIdx.y;
08          int tx = threadIdx.x; int ty = threadIdx.y;
09
10          // Identify the row and column of the P element to work on
11          int Row = by * TILE_WIDTH + ty;
12          int Col = bx * TILE_WIDTH + tx;
13
14          // Loop over the M and N tiles required to compute P element
15          float Pvalue = 0;
16          for (int ph = 0; ph < Width/TILE_WIDTH; ++ph) {
17
18              // Collaborative loading of M and N tiles into shared memory
19              Mds[ty][tx] = M[Row*Width + ph*TILE_WIDTH + tx];
20              Nds[ty][tx] = N[(ph*TILE_WIDTH + ty)*Width + Col];
21              __syncthreads();
22
23              for (int k = 0; k < TILE_WIDTH; ++k) {
24                  Pvalue += Mds[ty][k] * Nds[k][tx];
25              }
26              __syncthreads();
27
28          }
29          P[Row*Width + Col] = Pvalue;
30
31      }

```

**FIGURE 5.9**

---

A tiled matrix multiplication kernel using shared memory.